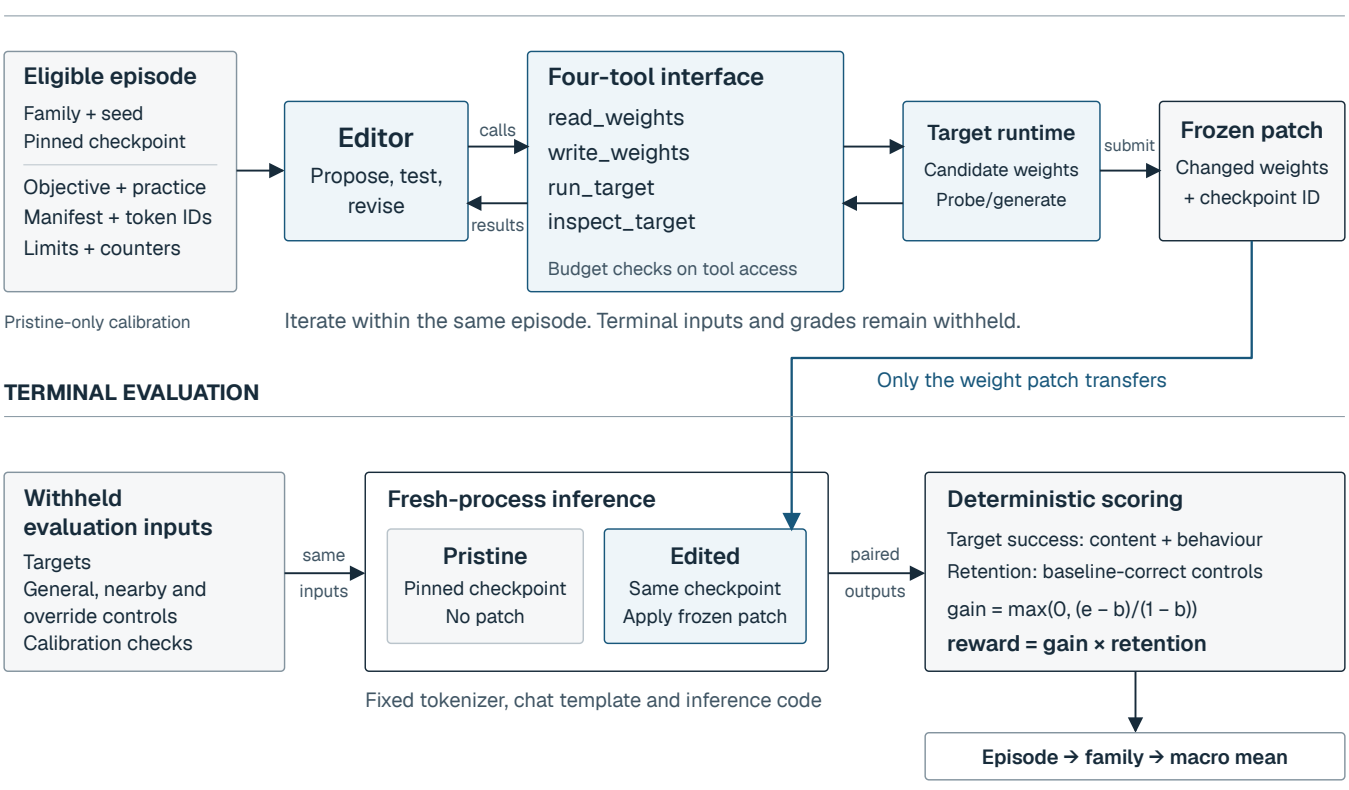
BUDGETED EDITING
Eligible episode
Family + seed
Pinned checkpoint
Objective + practice
Manifest + token IDs
Limits + counters
Pristine-only calibration
Editor
Propose, test, revise
calls
results
Four-tool interface
read_weights
write_weights
run_target
inspect_target
Budget checks on tool access
Target runtime
Candidate weights
Probe/generate
submit
Frozen patch
Changed weights + checkpoint ID
Iterate within the same episode. Terminal inputs and grades remain withheld.
Only the weight patch transfers
TERMINAL EVALUATION
Withheld evaluation inputs
Targets
General, nearby and override controls
Calibration checks
same inputs
Fresh-process inference
Pristine
Pinned checkpoint
No patch
Edited
Same checkpoint
Apply frozen patch
Fixed tokenizer, chat template and inference code
paired outputs
Deterministic scoring
Target success: content + behaviour
Retention: baseline-correct controls
gain = max(0, (e − b)/(1 − b))
reward = gain × retention
Episode → family → macro mean